

# MSIL MCP Server - Complete Architecture & Data Storage Guide

---

**Document Date:** January 31, 2026

**Version:** 1.0

**Status:** Complete Architecture Documentation

---

## Table of Contents

---

1. [Overview](#)
  2. [Admin API Addition Workflow](#)
  3. [Security Layers: Authentication, Authorization & Rate Limiting](#)
  4. [MCP Client Call Flow](#)
  5. [Data Storage Locations](#)
  6. [User & Role Management](#)
  7. [Tool Registry & Configuration](#)
  8. [API Gateway Configuration](#)
  9. [Security Policies & Rules](#)
  10. [RFP Requirements Mapping](#)
- 

## Overview

---

The MSIL MCP Server implements a **Model Context Protocol** (MCP) with enterprise-grade security for Maruti Suzuki India Limited's service booking platform.

### Key Architecture Components:

- **Zero-Code Tool Generation** from OpenAPI specs
  - **Multi-Layered Security** (Authentication → Authorization → Rate Limiting)
  - **Risk-Based Access Control** (Read/Write/Privileged tools)
  - **Real-Time Metrics & Audit Logging**
  - **Idempotent Write Operations**
  - **Dual API Gateway Support** (Mock API + MSIL APIM)
- 

## Admin API Addition Workflow

---

### Stage 1: OpenAPI Spec Upload

**Entry Point:** Admin UI → "Import OpenAPI" page

**Endpoint:** `POST /api/openapi-import/upload`

**Request:**

```
{
  "file": "<openapi_spec_yaml>",
  "category": "imported",
  "bundle_name": "customer_service_v1"
}
```

**Files Involved:**

- **Frontend:** admin-ui/src/pages/Import.tsx
- **Backend:** mcp-server/app/api/openapi\_import.py (lines 69-120)
- **Parser:** mcp-server/app/core/openapi/parser.py

**Processing:**

1. File is saved temporarily
2. OpenAPI YAML is parsed
3. All endpoints are extracted
4. For each endpoint, a **Tool object** is created with:
  - **name**: Auto-generated from endpoint path
  - **display\_name**: Formatted for UI
  - **description**: From OpenAPI spec
  - **api\_endpoint**: Full path
  - **http\_method**: GET/POST/PUT/DELETE/PATCH
  - **input\_schema**: From parameters + request body
  - **output\_schema**: From response schema
  - **risk\_level**: Auto-detected (read/write/privileged)

**Example Tool Generation:**

```

# Input (OpenAPI):
openapi: 3.0.0
paths:
  /api/customer/resolve:
    post:
      summary: Resolve Customer
      parameters:
        - name: mobile
          required: true
          schema:
            type: string
      responses:
        200:
          schema:
            type: object
            properties:
              customer_id: string
              name: string

# Output (Generated Tool):
Tool(
  id=uuid4(),
  name="resolve_customer",
  display_name="Resolve Customer",
  description="Resolve Customer",
  api_endpoint="/api/customer/resolve",
  http_method="POST",
  input_schema={
    "type": "object",
    "properties": {
      "mobile": {"type": "string"}
    },
    "required": ["mobile"]
  },
  output_schema={
    "type": "object",
    "properties": {
      "customer_id": {"type": "string"},
      "name": {"type": "string"}
    }
  },
  risk_level="read",
  rate_limit_tier="standard"
)

```

## Stage 2: Tool Preview & Selection

**Display:** Admin UI → "Import Preview" page

**Storage:** In-memory cache (temporary)

**File:** mcp-server/app/api/openapi\_import.py (lines 212-260)

**Data Cached in Memory:**

```
_specs_cache: Dict[str, Dict[str, Any]] = {
    "spec_abc123": {
        "name": "customer_service_api",
        "version": "1.0.0",
        "uploaded_at": "2026-01-31T10:30:45Z",
        "tools": [
            {ToolPreview object},
            {ToolPreview object},
            ...
        ],
        "status": "preview"
    }
}
```

**Admin Actions:**

- View all generated tools with metadata
  - Select/deselect tools to register
  - Modify tool properties (name, category, risk level)
  - Click "Register Selected Tools"
- 

**Stage 3: Approval & Registration****Endpoint:** POST /api/openapi-import/approve-tools**Request:**

```
{
    "spec_id": "spec_abc123",
    "tool_ids": ["resolve_customer", "resolve_vehicle",
    "get_nearby_dealers"],
    "category": "service_booking"
}
```

**Processing:**

1. Load selected tools from temporary cache
2. **Validate** each tool against schema
3. **Persist to Database** (PostgreSQL `tools` table)
4. **Reload tool registry** (in-memory cache invalidation)
5. Return confirmation

**Database Insertion:**

```
-- mcp-server/app/db/database.py
INSERT INTO tools (
    id, name, display_name, description, category,
    api_endpoint, http_method, input_schema, output_schema,
    risk_level, rate_limit_tier, is_active, created_at, updated_at
) VALUES (
    'uuid123', 'resolve_customer', 'Resolve Customer', 'Get customer
details...', 
    'service_booking', '/api/customer/resolve', 'POST',
    '{"type": "object", "properties": {...}}',
    '{"type": "object", "properties": {...}}',
    'read', 'standard', true, NOW(), NOW()
)
```

## Security Layers: Authentication, Authorization & Rate Limiting

### Layer 1: Authentication (JWT Token Validation)

File: [mcp-server/app/core/auth/oauth2\\_provider.py](#)

When Client Calls MCP:

```
POST /api/mcp
Headers:
    Authorization: Bearer eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9...
    X-API-Key: msil-mcp-dev-key-2026
    X-Correlation-ID: req-12345
```

Authentication Handler:

```
# Line 95-120: get_current_user()
async def get_current_user(
    credentials: HTTPAuthorizationCredentials = Depends(security)
) -> UserInfo:
    """
    Extract and validate JWT token
    Returns: UserInfo with user_id, email, roles
    """
    token = credentials.credentials # Bearer token
    payload = jwt_handler.decode_token(token) # Validate signature

    # Extracted from JWT claims:
    user_info = UserInfo(
        user_id=payload["user_id"],
        email=payload["email"],
        name=payload["name"],
        roles=payload["roles"], # ["developer", "operator"]
        is_active=payload["is_active"]
    )
    return user_info
```

**JWT Token Structure:**

```
# Encoded JWT contains:
{
    "header": {
        "alg": "HS256",
        "typ": "JWT"
    },
    "payload": {
        "user_id": "usr_dev_001",
        "email": "developer@msil.com",
        "name": "MSIL Developer",
        "roles": ["developer", "operator"],
        "is_active": true,
        "exp": 1737986400, # Expires in 60 minutes
        "iat": 1737982800
    },
    "signature": "HMACSHA256(header + payload, JWT_SECRET)"
}
```

**JWT Configuration (app/config.py):**

```
JWT_SECRET: str = "msil-mcp-jwt-secret-key-change-in-production-2026"
JWT_ALGORITHM: str = "HS256"
JWT_ACCESS_TOKEN_EXPIRE_MINUTES: int = 60
JWT_REFRESH_TOKEN_EXPIRE_DAYS: int = 7
```

**Demo Mode (MVP):**

```
# If DEMO_MODE=True, hardcoded users bypass actual JWT validation:
self.users = {
    "admin@msil.com": {
        "user_id": "usr_admin_001",
        "password_hash": "admin123",
        "roles": ["admin", "developer", "operator"]
    },
    "developer@msil.com": {
        "user_id": "usr_dev_001",
        "password_hash": "dev123",
        "roles": ["developer", "operator"]
    },
    "operator@msil.com": {
        "user_id": "usr_op_001",
        "password_hash": "op123",
        "roles": ["operator"]
    }
}
```

## Layer 2: Authorization & Risk-Based Policy

### Files:

- mcp-server/app/core/policy/engine.py
- mcp-server/app/core/policy/risk\_policy.py
- mcp-server/app/core/policy/models.py

### Policy Evaluation Context:

```
# Context built from request:
context = {
    "action": "invoke",           # discover, invoke, admin
    "resource": "resolve_customer", # tool name
    "user_id": "usr_dev_001",
    "roles": ["developer", "operator"],
    "tool_risk_level": "read",    # from Tool.risk_level
    "is_elevated": False,         # PIM elevation status
    "ip_address": "192.168.1.100",
    "timestamp": "2026-01-31T10:30:45Z"
}
```

### Risk Policy Rules:

```
# File: mcp-server/app/core/policy/risk_policy.py

class RiskPolicy:
    """Defines access control for each risk level"""
    risk_level: str          # "read", "write", "privileged"
    min_role: str             # Minimum role required
    requires_elevation: bool # Needs PIM/PAM elevation
    requires_approval: bool  # Needs manager approval
    rate_limit_tier: str     # Rate Limiting tier
    max_concurrency: int    # Max concurrent executions
    pii_policy: str          # PII handling policy

# Initialized policies:
POLICIES = {
    "read": RiskPolicy(
        risk_level="read",
        min_role="operator",
        requires_elevation=False,
        requires_approval=False,
        rate_limit_tier="permissive", # 100 req/min
        max_concurrency=50,
        pii_policy="redact"
    ),
    "write": RiskPolicy(
        risk_level="write",
        min_role="developer",
        requires_elevation=False,
        requires_approval=False,
        rate_limit_tier="standard",   # 50 req/min
        max_concurrency=20,
        pii_policy="mask"
    ),
    "privileged": RiskPolicy(
        risk_level="privileged",
        min_role="admin",
        requires_elevation=True,      # Requires PIM approval
        requires_approval=True,       # Requires manager approval
        rate_limit_tier="strict",     # 10 req/min
        max_concurrency=5,
        pii_policy="redact_sensitive"
    )
}
```

### Access Decision Logic:

```
# File: mcp-server/app/core/policy/risk_policy.py (Line 93-150)

def evaluate_access(
    self,
    tool_risk_level: str,
    user_role: str,
    is_elevated: bool = False
) -> Dict[str, Any]:
    """
    Evaluate if user can execute tool
    """

    policy = self.get_policy(tool_risk_level)

    # Role hierarchy check
    role_hierarchy = {
        "user": 0,
        "operator": 1,
        "developer": 2,
        "admin": 3
    }

    user_level = role_hierarchy.get(user_role, -1)
    required_level = role_hierarchy.get(policy.min_role, 999)
    has_role = user_level >= required_level

    # Elevation check
    needs_elevation = policy.requires_elevation and not is_elevated

    # Final decision
    allowed = has_role and not needs_elevation

    return {
        "allowed": allowed,
        "has_required_role": has_role,
        "requires_elevation": policy.requires_elevation,
        "is_elevated": is_elevated,
        "needs_elevation": needs_elevation,
        "requires_approval": policy.requires_approval,
        "rate_limit_tier": policy.rate_limit_tier,
        "max_concurrency": policy.max_concurrency,
        "reason": "Access denied: User 'operator' cannot execute 'privileged' tool"
    }
```

### Example Scenarios:

User	Role	Tool	Risk	Elevation	Result	Reason
user1	operator	resolve_customer	read	No	<span style="color: green;">✓</span> ALLOW	Role meets requirement
user2	operator	delete_booking	privileged	No	<span style="color: red;">✗</span> DENY	Requires admin role

User	Role	Tool	Risk	Elevation	Result	Reason
user3	admin	delete_booking	privileged	No	X DENY	Elevation required
user3	admin	delete_booking	privileged	Yes	✓ ALLOW	All checks pass

## Layer 3: Rate Limiting (Token Bucket)

**File:** mcp-server/app/core/cache/rate\_limiter.py

**Storage:** Redis (configured in app/config.py)

**Rate Limit Tiers:**

```
RATE_LIMITS = {
    "permissive": {
        "requests_per_minute": 100,
        "burst_size": 20,           # Allow 20 simultaneous requests
        "window": 60                # 1 minute window
    },
    "standard": {
        "requests_per_minute": 50,
        "burst_size": 10,
        "window": 60
    },
    "strict": {
        "requests_per_minute": 10,
        "burst_size": 2,
        "window": 60
    }
}
```

**Rate Limit Check Process:**

```
# File: mcp-server/app/core/cache/rate_limiter.py (Line 42-85)
```

```
async def check_rate_limit(
    self,
    key: str,           # e.g., "user:usr_dev_001" or
"tool:resolve_customer"
    limit: int,         # e.g., 100
    window: int = 60    # seconds
) -> RateLimitInfo:
    """
    Token bucket algorithm:
    - Each second, add 'limit/window' tokens
    - Each request consumes 1 token
    - If tokens available: allowed
    - If tokens exhausted: reject with retry_after
    """
    cache_key = f"ratelimit:{key}"
    current_time = int(time.time())

    # Get current count from Redis
    count = await redis.increment(cache_key, 1, window)

    if count <= limit:
        # Request allowed
        return RateLimitInfo(
            allowed=True,
            remaining=limit - count,
            reset_at=current_time + window
        )
    else:
        # Rate limit exceeded
        retry_after = window - (current_time % window)
        return RateLimitInfo(
            allowed=False,
            remaining=0,
            reset_at=current_time + retry_after,
            retry_after=retry_after
        )
```

### Per-User & Per-Tool Limits:

```

# Check user limit
user_limit = await rate_limiter.check_user_rate_limit(
    user_id="usr_dev_001",
    limit=100, # 100 requests/minute for this user
    window=60
)
if not user_limit.allowed:
    return 429 Too Many Requests

# Check tool limit
tool_limit = await rate_limiter.check_tool_rate_limit(
    tool_name="resolve_customer",
    limit=50, # 50 requests/minute for this tool
    window=60
)
if not tool_limit.allowed:
    return 429 Too Many Requests

```

**Response Headers:**

```

HTTP/1.1 200 OK
X-RateLimit-Limit: 100
X-RateLimit-Remaining: 87
X-RateLimit-Reset: 1737980000

# If rate limited:
HTTP/1.1 429 Too Many Requests
Retry-After: 45
X-RateLimit-Remaining: 0
X-RateLimit-Reset: 1737980045

```

**MCP Client Call Flow****Complete End-to-End Journey****1. CLIENT INITIATES REQUEST**

```

POST /api/mcp
Headers:
    Content-Type: application/json
    Authorization: Bearer eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9...
    X-API-Key: msil-mcp-dev-key-2026
    X-Correlation-ID: req-12345
    Idempotency-Key: idempotent-abc123 # For write operations

Body:
{

```

```

"jsonrpc": "2.0",
"id": "1",
"method": "tools/call",
"params": {
    "name": "resolve_customer",
    "arguments": {
        "mobile": "9876543210"
    }
}

```



## 2. AUTHENTICATION (oauth2\_provider.py)

- ✓ Extract JWT from Authorization header
- ✓ Validate signature using JWT\_SECRET (HS256)
- ✓ Check token expiration
- ✓ Extract claims: user\_id, email, roles

UserInfo Extracted:

```
{
    "user_id": "usr_dev_001",
    "email": "developer@msil.com",
    "name": "MSIL Developer",
    "roles": ["developer", "operator"],
    "is_active": true
}
```

- ✓ Authentication passed → Continue
- ✗ Auth failed → Return 401 Unauthorized



## 3. RATE LIMIT CHECK - PER USER (rate\_limiter.py)

- ✓ Get user\_id: "usr\_dev\_001"
- ✓ Get user's rate limit tier from policy: "standard"
- ✓ Check Redis: key="ratelimit:user:usr\_dev\_001"
- ✓ Decrement token count (from 50/minute)
  
- ✓ Tokens available → Continue
- ✗ Limit exceeded → Return 429 Too Many Requests  
Retry-After: 45 seconds



## 4. PARSE & ROUTE REQUEST (app/api/mcp.py)

2026-01-31

- ✓ Parse JSON-RPC: method="tools/call"
- ✓ Generate/use correlation\_id: "req-12345"
- ✓ Route to handle\_tools\_call()



#### 5. TOOL DISCOVERY (app/core/tools/registry.py)

- ✓ Get tool\_name: "resolve\_customer"
- ✓ Look up in registry (in-memory cache)

Tool object returned:

```
{  
    "id": "uuid123",  
    "name": "resolve_customer",  
    "display_name": "Resolve Customer",  
    "api_endpoint": "/api/customer/resolve",  
    "http_method": "POST",  
    "input_schema": {...},  
    "output_schema": {...},  
    "risk_level": "read",  
    "rate_limit_tier": "standard",  
    "requires_approval": false,  
    "requires_elevation": false  
}
```

- ✓ Tool found → Continue
- ✗ Tool not found → Return -32602 Invalid params



#### 6. INPUT VALIDATION (Pydantic + JSON Schema)

- ✓ Validate arguments against tool.input\_schema
  - ✓ Check required fields: "mobile" present? Yes
  - ✓ Check field types: string? Yes
  - ✓ Check payload size: < 1MB? Yes
  - ✓ Sanitize inputs (no injection patterns)
- 
- ✓ Validation passed → Continue
  - ✗ Validation failed → Return -32602 Invalid params with details



#### 7. AUTHORIZATION & RISK POLICY (engine.py)

Build policy context:

```
{  
    "action": "invoke",  
    "resource": "resolve_customer"
```

2026-01-31

```
resource . resolve_customer ,  
"user_id": "usr_dev_001",  
"roles": ["developer", "operator"],  
"tool_risk_level": "read",  
"is_elevated": false  
}
```

Policy evaluation:

- ✓ Tool risk="read" → min\_role="operator"
- ✓ User roles include "developer"  (exceeds requirement)
- ✓ Elevation NOT required
- ✓ Approval NOT required
  
- ✓ Access allowed → Continue
- X Access denied → Return -32001 Unauthorized  
"Reason: User 'operator' cannot execute 'privileged' tool"



#### 8. RATE LIMIT CHECK - PER TOOL (rate\_limiter.py)

- ✓ Get tool rate limit: "standard" tier = 50 req/min
- ✓ Check Redis: key="ratelimit:tool:resolve\_customer"
- ✓ Decrement token count
  
- ✓ Tokens available → Continue
- X Limit exceeded → Return 429 Too Many Requests



#### 9. IDEMPOTENCY CHECK (For Write Operations)

Only for POST/PUT/PATCH/DELETE methods:

- ✓ Get idempotency\_key from header: "idempotent-abc123"
- ✓ Look up in Redis: "idempotency:idempotent-abc123"
  
- ✓ Cache HIT: Return cached response (no re-execution)
- ✓ Cache MISS: Continue, will cache result after execution



#### 10. TOOL EXECUTION (app/core/tools/executor.py)

Determine API gateway:

```
If settings.API_GATEWAY_MODE == "mock":  
    base_url = "http://localhost:8080"  
Else if settings.API_GATEWAY_MODE == "msil_apim":  
    base_url = "https://apim-dev.marutisuzuki.com"
```

```
Build HTTP request:
{
    "method": "POST",
    "url": "http://localhost:8080/api/customer/resolve",
    "headers": {
        "Content-Type": "application/json",
        "X-API-Key": "msil-mcp-dev-key-2026",
        "X-Correlation-ID": "req-12345",
        "X-User-Context": "usr_dev_001"
    },
    "json": {
        "mobile": "9876543210"
    },
    "timeout": 30.0
}
```

Execute with retry logic:

```
Max retries: 3
Backoff: 100ms, 200ms, 400ms exponential
```

Response received:

```
{
    "customer_id": "C123456",
    "name": "Rajesh Kumar",
    "phone": "9876543210",
    "vehicle_count": 2
}
```

```
✓ Success (200-299) → Extract response
✗ Timeout → Return -32603 Internal error
✗ 4xx → Return -32602 Invalid params
✗ 5xx → Return -32603 Internal error
```



### | 11. METRICS & AUDIT LOGGING (metrics/collector.py + audit.py) |

Record metrics:

- Tool: "resolve\_customer"
- User: "usr\_dev\_001"
- Status: "success"
- Duration: 245 ms
- Timestamp: 2026-01-31T10:30:45Z
- Request size: 85 bytes
- Response size: 156 bytes

Record audit log:

- User: "developer@msil.com"
- Action: "tool\_executed"
- Tool: "resolve\_customer"
- Input: {"mobile": "9876...3210"} # PII masked
- Output: "success"
- Correlation ID: "req-12345"

- Risk level: "read"



## 12. RESPONSE FORMATTING (JSON-RPC 2.0)

HTTP/1.1 200 OK

Content-Type: application/json  
 X-RateLimit-Remaining: 86  
 X-RateLimit-Reset: 1737980045  
 X-Correlation-ID: req-12345

```
{
  "jsonrpc": "2.0",
  "id": "1",
  "result": {
    "content": [
      {
        "type": "text",
        "text": "Customer resolved successfully"
      },
      {
        "type": "json",
        "data": {
          "customer_id": "C123456",
          "name": "Rajesh Kumar",
          "phone": "9876543210",
          "vehicle_count": 2
        }
      }
    ],
    "isError": false
  }
}
```

## Data Storage Locations

### 1. User & Authentication Data

**Storage Type:** In-Memory (Demo) + Database (Production)

**File Locations:**

- **Demo Users:** [mcp-server/app/core/auth/oauth2\\_provider.py](#) (lines 35-60)
- **JWT Handler:** [mcp-server/app/core/auth/jwt\\_handler.py](#)

**Demo Users:**

```

self.users = {
    "admin@msil.com": {
        "user_id": "usr_admin_001",
        "password_hash": "admin123",
        "roles": ["admin", "developer", "operator"],
        "is_active": True
    },
    "developer@msil.com": {
        "user_id": "usr_dev_001",
        "password_hash": "dev123",
        "roles": ["developer", "operator"],
        "is_active": True
    },
    "operator@msil.com": {
        "user_id": "usr_op_001",
        "password_hash": "op123",
        "roles": ["operator"],
        "is_active": True
    }
}

```

**Database Schema** (PostgreSQL):

```

-- Table: users
CREATE TABLE users (
    id UUID PRIMARY KEY,
    email VARCHAR(255) UNIQUE NOT NULL,
    name VARCHAR(255),
    password_hash VARCHAR(255),
    roles TEXT[] DEFAULT ARRAY['user'], -- Array of role strings
    is_active BOOLEAN DEFAULT TRUE,
    created_at TIMESTAMP DEFAULT NOW(),
    updated_at TIMESTAMP DEFAULT NOW()
);

-- Sample data:
INSERT INTO users VALUES (
    'usr_dev_001',
    'developer@msil.com',
    'MSIL Developer',
    'bcrypt_hash_...',
    ARRAY['developer', 'operator'],
    TRUE,
    NOW(),
    NOW()
);

```

**Configuration:** [mcp-server/app/config.py](#) (lines 63-72)

```

OAUTH2_ENABLED: bool = True
JWT_SECRET: str = "msil-mcp-jwt-secret-key-change-in-production-2026"
JWT_ALGORITHM: str = "HS256"
JWT_ACCESS_TOKEN_EXPIRE_MINUTES: int = 60
JWT_REFRESH_TOKEN_EXPIRE_DAYS: int = 7

```

## 2. Tools & Tool Registry

**Storage Type:** Database (Primary) + In-Memory Cache (Performance)

**File Locations:**

- **Registry:** mcp-server/app/core/tools/registry.py
- **Database:** mcp-server/app/db/database.py
- **Repositories:** mcp-server/app/db/repositories.py

**Database Schema:**

```

-- Table: tools
CREATE TABLE tools (
    id UUID PRIMARY KEY,
    name VARCHAR(255) UNIQUE NOT NULL,
    display_name VARCHAR(255),
    description TEXT,
    category VARCHAR(100),
    api_endpoint VARCHAR(500),
    http_method VARCHAR(10),    -- GET, POST, PUT, DELETE, PATCH
    input_schema JSONB,        -- JSON schema for inputs
    output_schema JSONB,       -- JSON schema for outputs
    headers JSONB,             -- Default headers
    auth_type VARCHAR(50),     -- none, basic, bearer, api_key
    is_active BOOLEAN DEFAULT TRUE,
    version VARCHAR(20) DEFAULT '1.0.0',
    risk_level VARCHAR(50) DEFAULT 'read', -- read, write, privileged
    requires_elevation BOOLEAN DEFAULT FALSE,
    requires_approval BOOLEAN DEFAULT FALSE,
    max_concurrent_executions INTEGER DEFAULT 10,
    rate_limit_tier VARCHAR(50) DEFAULT 'standard',
    created_at TIMESTAMP DEFAULT NOW(),
    updated_at TIMESTAMP DEFAULT NOW()
);

-- Indexes for performance:
CREATE INDEX idx_tools_name ON tools(name);
CREATE INDEX idx_tools_category ON tools(category);
CREATE INDEX idx_tools_is_active ON tools(is_active);

```

**In-Memory Cache (Registry):**

```
# File: mcp-server/app/core/tools/registry.py (Lines 42-50)
class ToolRegistry:
    def __init__(self):
        self._tools: Dict[str, Tool] = {} # Key: tool_name
        self._loaded = False

    # Loaded on startup:
    _tools = {
        "resolve_customer": Tool(...),
        "resolve_vehicle": Tool(...),
        "get_nearby_dealers": Tool(...),
        ...
    }
```

### Tool Object Structure:

```
@dataclass
class Tool:
    id: uuid.UUID
    name: str
    display_name: str
    description: str
    category: str
    api_endpoint: str
    http_method: str
    input_schema: Dict[str, Any]
    output_schema: Optional[Dict[str, Any]]
    headers: Dict[str, str]
    auth_type: str
    is_active: bool
    version: str
    risk_level: str
    requires_elevation: bool
    requires_approval: bool
    max_concurrent_executions: int
    rate_limit_tier: str
    strict
    created_at: Optional[datetime]
    updated_at: Optional[datetime]
```

### Loading from Database:

```
# File: mcp-server/app/core/tools/registry.py (Lines 53-85)
async def _load_from_db(self):
    async with get_db_session() as session:
        result = await session.execute(
            text("SELECT * FROM tools WHERE is_active = true"))
    rows = result.fetchall()
    for row in rows:
        tool = Tool(
            id=row.id,
            name=row.name,
            display_name=row.display_name,
            # ... map all fields
        )
        self._tools[tool.name] = tool
```

---

### 3. Tool Execution History & Metrics

**Storage Type:** Time-Series Database (TimescaleDB) or Prometheus

**File Locations:**

- **Collector:** [mcp-server/app/core/metrics/collector.py](#)
- **Analytics API:** [mcp-server/app/api/analytics.py](#)

**Database Schema:**

```
-- Table: tool_executions (time-series)
CREATE TABLE tool_executions (
    id UUID PRIMARY KEY,
    correlation_id VARCHAR(255),
    tool_name VARCHAR(255),
    user_id VARCHAR(255),
    status VARCHAR(50),           -- success, failure, timeout
    error_code INTEGER,
    error_message TEXT,
    duration_ms INTEGER,         -- Execution time in milliseconds
    request_size_bytes INTEGER,
    response_size_bytes INTEGER,
    input_data JSONB,            -- Sanitized input
    risk_level VARCHAR(50),
    created_at TIMESTAMP DEFAULT NOW()
);

-- Create index on time for fast range queries:
CREATE INDEX idx_executions_time ON tool_executions(created_at DESC);
CREATE INDEX idx_executions_tool ON tool_executions(tool_name);
CREATE INDEX idx_executions_user ON tool_executions(user_id);

-- Example queries:
SELECT tool_name, COUNT(*) as call_count, AVG(duration_ms) as avg_duration
FROM tool_executions
WHERE created_at >= NOW() - INTERVAL '1 day'
GROUP BY tool_name
ORDER BY call_count DESC;
```

## Metrics Recorded:

```
@dataclass
class ToolMetrics:
    tool_name: str
    user_id: str
    status: str           # success, failure
    duration_ms: int
    timestamp: datetime
    request_size_bytes: int
    response_size_bytes: int
    error_code: Optional[int]
    error_message: Optional[str]
    risk_level: str
```

---

## 4. Audit Logs

**Storage Type:** Immutable Log Storage (PostgreSQL or Elasticsearch)

**File Locations:**

- **Audit Service:** [mcp-server/app/core/audit/service.py](#)
- **PII Masker:** [mcp-server/app/core/audit/pii\\_masker.py](#)

### Database Schema:

```
-- Table: audit_logs (immutable)
CREATE TABLE audit_logs (
    id UUID PRIMARY KEY,
    timestamp TIMESTAMP DEFAULT NOW(),
    correlation_id VARCHAR(255),
    user_id VARCHAR(255),
    user_email VARCHAR(255),
    action VARCHAR(100),          -- tool_executed, tool_registered, etc.
    resource_type VARCHAR(100),   -- tool, spec, policy
    resource_id VARCHAR(255),
    details JSONB,               -- Action details (PII masked)
    status VARCHAR(50),          -- success, failure, denied
    reason_code VARCHAR(100),     -- For denials
    ip_address VARCHAR(50),
    user_agent TEXT
);

-- Example audit log:
{
  "id": "audit-123",
  "timestamp": "2026-01-31T10:30:45Z",
  "correlation_id": "req-12345",
  "user_id": "usr_dev_001",
  "user_email": "developer@msil.com",
  "action": "tool_executed",
  "resource_type": "tool",
  "resource_id": "resolve_customer",
  "details": {
    "input": {"mobile": "9876...3210"},      # PII masked
    "output": "success",
    "duration_ms": 245,
    "risk_level": "read"
  },
  "status": "success",
  "ip_address": "192.168.1.100"
}
```

### PII Masking Rules:

```
# File: mcp-server/app/core/audit/pii_masker.py

PII_PATTERNS = {
    "phone": r"\d{10}",                      # 10-digit phone
    "email": r"[\w\.-]+@[\\w\.-]+\.",      # Email addresses
    "pan": r"[A-Z]{5}[0-9]{4}[A-Z]",        # PAN numbers
    "aadhaar": r"\d{12}",                   # Aadhaar (12 digits)
    "vehicle_reg": r"[A-Z]{2}\d{2}[A-Z]{2}\d{4}" # Registration plate
}

# Masking applied:
"9876543210" → "9876...3210"
"dev@example.com" → "dev@*****.com"
```

## 5. Policy & Configuration Data

**Storage Type:** Database + In-Memory Cache

**File Locations:**

- **Policy Engine:** mcp-server/app/core/policy/engine.py
- **Risk Policy:** mcp-server/app/core/policy/risk\_policy.py
- **Configuration:** mcp-server/app/config.py

**Database Schema:**

```
-- Table: policies
CREATE TABLE policies (
    id UUID PRIMARY KEY,
    policy_type VARCHAR(100), -- rbac, rate_limit, risk
    policy_name VARCHAR(255),
    content JSONB,           -- Policy rules
    version INTEGER DEFAULT 1,
    is_active BOOLEAN DEFAULT TRUE,
    created_at TIMESTAMP DEFAULT NOW(),
    updated_at TIMESTAMP DEFAULT NOW()
);

-- Example RBAC policy:
{
    "policy_type": "rbac",
    "policy_name": "default_roles",
    "content": {
        "admin": ["tools:read", "tools:write", "policies:write"],
        "developer": ["tools:read", "tools:write"],
        "operator": ["tools:read"],
        "user": ["tools:invoke"]
    }
}

-- Example rate limit policy:
{
    "policy_type": "rate_limit",
    "policy_name": "standard_limits",
    "content": {
        "permissive": {"requests_per_minute": 100, "burst_size": 20},
        "standard": {"requests_per_minute": 50, "burst_size": 10},
        "strict": {"requests_per_minute": 10, "burst_size": 2}
    }
}
```

**Configuration File** ([app/config.py](#)):

```

# Lines 63-100: Security & Policy Configuration

# OAuth2/JWT
OAUTH2_ENABLED: bool = True
JWT_SECRET: str = "msil-mcp-jwt-secret-key-..."
JWT_ALGORITHM: str = "HS256"
JWT_ACCESS_TOKEN_EXPIRE_MINUTES: int = 60

# API Gateway Mode
API_GATEWAY_MODE: str = "mock" # or "msil_apim"
MOCK_API_BASE_URL: str = "http://localhost:8080"
MSIL_APIM_BASE_URL: str = "https://apim-dev.marutisuzuki.com"
MSIL_APIM_SUBSCRIPTION_KEY: Optional[str] = None

# Rate Limiting
REDIS_ENABLED: bool = True
REDIS_URL: str = "redis://localhost:6379/0"

# PII & Security
PII_MASKING_ENABLED: bool = True
IDEMPOTENCY_ENABLED: bool = True
IDEMPOTENCY_REQUIRED: bool = False

# Features
DEMO_MODE: bool = True # Bypass auth in MVP
OIDC_ENABLED: bool = False # For production

```

## 6. Cache & Session Storage

**Storage Type:** Redis

**File Locations:**

- **Cache Service:** mcp-server/app/core/cache/service.py
- **Rate Limiter:** mcp-server/app/core/cache/rate\_limiter.py
- **Idempotency Store:** mcp-server/app/core/idempotency/store.py

**Redis Keys Structure:**

```

# Rate limiting
ratelimit:user:usr_dev_001 → 45 (remaining tokens)
ratelimit:tool:resolve_customer → 38

# Idempotency
idempotency:idempotent-abc123 → {response_json} # TTL: 24 hours

# Session cache
session:req-12345 → {session_data} # TTL: 1 hour

# Tool registry cache
tools:list → [tool1, tool2, ...] # TTL: 5 minutes
tool:resolve_customer → {tool_data} # TTL: 5 minutes

# OPA policies (if enabled)
policy:rbac → {policy_json} # TTL: 1 hour

```

### Redis Configuration (app/config.py):

```

REDIS_URL: str = "redis://localhost:6379/0"
REDIS_ENABLED: bool = True
CACHE_TTL: int = 300      # 5 minutes for tool cache
CACHE_MAX_SIZE: int = 1000 # Max cache entries

```

## 7. OpenAPI Specs & Tool Definitions

**Storage Type:** File System + Database

**File Locations:**

- **Sample Specs:** sample-apis/customer-service-api.yaml
- **Upload Handler:** mcp-server/app/api/openapi\_import.py
- **Parser:** mcp-server/app/core/openapi/parser.py

**File Storage (Temporary):**

```

mcp-server/
  └── uploads/
    ├── spec_abc123.yaml
    ├── spec_def456.yaml
    └── ...

```

**Database Schema (Persistent):**

```
-- Table: openapi_specs
CREATE TABLE openapi_specs (
    id UUID PRIMARY KEY,
    name VARCHAR(255),
    version VARCHAR(50),
    description TEXT,
    spec_content JSONB,          -- Full OpenAPI spec
    category VARCHAR(100),
    bundle_name VARCHAR(255),
    tools_count INTEGER,
    status VARCHAR(50),          -- uploaded, parsed, approved
    uploaded_by VARCHAR(255),
    uploaded_at TIMESTAMP,
    approved_at TIMESTAMP,
    is_active BOOLEAN DEFAULT TRUE,
    created_at TIMESTAMP DEFAULT NOW()
);

-- Example:
{
    "id": "spec_abc123",
    "name": "customer_service_api",
    "version": "1.0.0",
    "category": "service_booking",
    "bundle_name": "customer_service_v1",
    "tools_count": 11,
    "status": "approved",
    "uploaded_by": "admin@msil.com",
    "uploaded_at": "2026-01-31T09:15:00Z",
    "approved_at": "2026-01-31T09:45:00Z"
}
```

### Sample OpenAPI Spec:

```
# sample-apis/customer-service-api.yaml
openapi: 3.0.0
info:
  title: Customer Service API
  version: 1.0.0
  description: Maruti Suzuki Service Booking API
paths:
  /api/customer/resolve:
    post:
      summary: Resolve Customer
      operationId: resolve_customer
      parameters:
        - name: mobile
          in: query
          required: true
          schema:
            type: string
      responses:
        200:
          description: Success
          content:
            application/json:
              schema:
                type: object
                properties:
                  customer_id:
                    type: string
                  name:
                    type: string
```

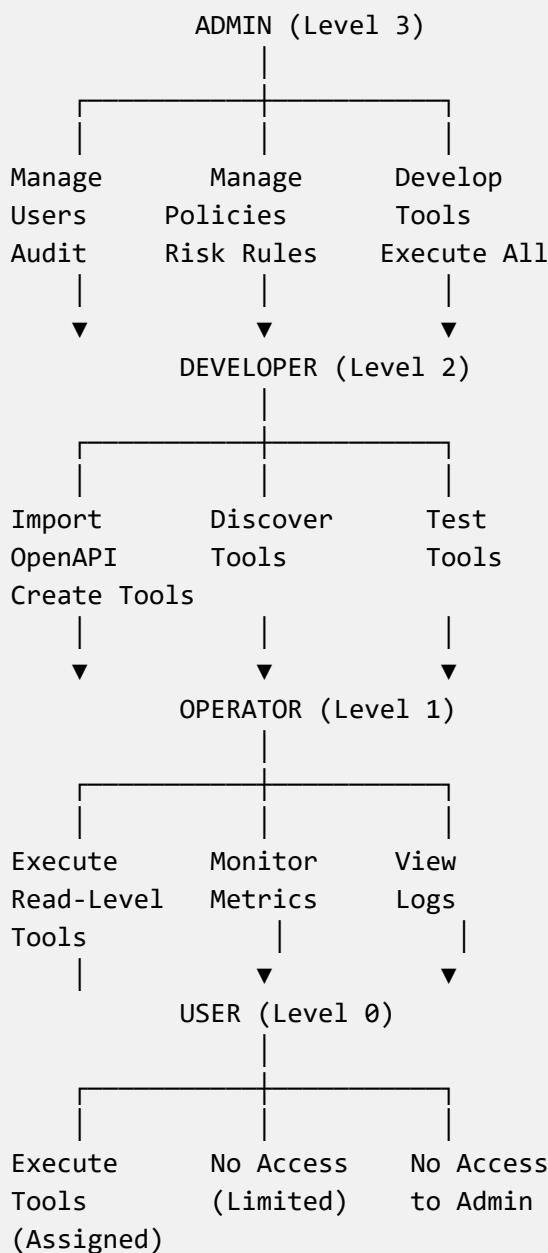
---

## User & Role Management

---

### Role Hierarchy

### USER ROLES & PERMISSIONS HIERARCHY



## Role Permissions Matrix

Permission	Admin	Developer	Operator	User
<b>Tool Discovery</b>	<input checked="" type="checkbox"/> All	<input checked="" type="checkbox"/> All	<input checked="" type="checkbox"/> All	<input checked="" type="checkbox"/> Assigned
<b>Tool Execution</b>	<input checked="" type="checkbox"/> All (Privileged)	<input checked="" type="checkbox"/> Write Level	<input checked="" type="checkbox"/> Read Level	<input checked="" type="checkbox"/> Limited
<b>OpenAPI Import</b>	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No	<input type="checkbox"/> No
<b>Tool Registration</b>	<input checked="" type="checkbox"/> Approve	<input checked="" type="checkbox"/> Create	<input type="checkbox"/> No	<input type="checkbox"/> No
<b>Policy Management</b>	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No	<input type="checkbox"/> No	<input type="checkbox"/> No
<b>User Management</b>	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No	<input type="checkbox"/> No	<input type="checkbox"/> No

Permission	Admin	Developer	Operator	User
<b>Audit Logs</b>	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Own Only	<input checked="" type="checkbox"/> Read	<input checked="" type="checkbox"/> No
<b>Metrics</b>	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> No
<b>PIM Elevation</b>	<input checked="" type="checkbox"/> Required	<input checked="" type="checkbox"/> No	<input checked="" type="checkbox"/> No	<input checked="" type="checkbox"/> No

## Demo Users

# File: mcp-server/app/core/auth/oauth2\_provider.py (lines 35-60)

```
DEMO_USERS = {
    "admin@msil.com": {
        "user_id": "usr_admin_001",
        "email": "admin@msil.com",
        "password_hash": "admin123",
        "name": "MSIL Admin",
        "roles": ["admin", "developer", "operator"],
        "is_active": True
    },
    "developer@msil.com": {
        "user_id": "usr_dev_001",
        "email": "developer@msil.com",
        "password_hash": "dev123",
        "name": "MSIL Developer",
        "roles": ["developer", "operator"],
        "is_active": True
    },
    "operator@msil.com": {
        "user_id": "usr_op_001",
        "email": "operator@msil.com",
        "password_hash": "op123",
        "name": "MSIL Operator",
        "roles": ["operator"],
        "is_active": True
    }
}
```

## Tool Registry & Configuration

### Tool Registry State

Location: [mcp-server/app/core/tools/registry.py](#)

Registry Lifecycle:

APPLICATION STARTUP

- Create ToolRegistry instance
- Initialize empty \_tools dict
- \_loaded = False



FIRST REQUEST (Lazy Loading)

- Call \_ensure\_loaded()
- Check if \_loaded = True  
NO → Call \_load\_from\_db()
- Query: SELECT \* FROM tools WHERE is\_active = true
- For each row, create Tool object
  - Populate \_tools dict: {"tool\_name": Tool(...)}
  - Set \_loaded = True



NORMAL OPERATIONS

- get\_tool("resolve\_customer")
  - Return \_tools["resolve\_customer"]
- list\_tools(category="service\_booking")
  - Filter \_tools by category
- count\_tools()
  - Return len(\_tools)



TOOL REGISTRATION (After OpenAPI Import)

- INSERT into tools table
- Call registry.reload() or restart app
- Set \_loaded = False (invalidate cache)
- Next request will re-load from database
  - New tools now available

## Tool Metadata

```

# Tool object with all configuration:

Tool(
    # Identity
    id=UUID("550e8400-e29b-41d4-a716-446655440000"),
    name="resolve_customer",
    display_name="Resolve Customer",
    description="Get customer details from phone number",

    # Routing
    api_endpoint="/api/customer/resolve",
    http_method="POST",
    category="service_booking",

    # Schema
    input_schema={
        "type": "object",
        "properties": {
            "mobile": {
                "type": "string",
                "pattern": "^[0-9]{10}$",
                "description": "10-digit mobile number"
            },
            "required": ["mobile"],
            "additionalProperties": False
        },
        output_schema={
            "type": "object",
            "properties": {
                "customer_id": {"type": "string"},
                "name": {"type": "string"},
                "phone": {"type": "string"},
                "vehicle_count": {"type": "integer"}
            }
        },
    },
    # Authentication
    auth_type="api_key",
    headers={"X-API-Key": "msil-mcp-dev-key-2026"},

    # Access Control
    risk_level="read",
    requires_elevation=False,
    requires_approval=False,
    max_concurrent_executions=50,
    rate_limit_tier="standard",

    # Status
    is_active=True,
    version="1.0.0",
    created_at=datetime(2026, 1, 31, 9, 30, 0),
    updated_at=datetime(2026, 1, 31, 9, 30, 0)
)

```

# API Gateway Configuration

## API Gateway Modes

File: `mcp-server/app/config.py` (lines 41-54)

### Mode 1: Mock API (Development)

```
API_GATEWAY_MODE: str = "mock"
MOCK_API_BASE_URL: str = "http://localhost:8080"

# When tool is executed:
# Tool endpoint: /api/customer/resolve
# Actual URL: http://localhost:8080/api/customer/resolve
```

### Mock API Server:

- **Location:** `mock-api/app/main.py`
- **Port:** 8080
- **Purpose:** Simulate Maruti Suzuki backend APIs for development/testing

### Mock API Endpoints:

```
# mock-api/app/routers/customer.py
POST /api/customer/resolve
- Input: {"mobile": "9876543210"}
- Output: {"customer_id": "C123456", "name": "Rajesh Kumar", ...}

POST /api/vehicle/resolve
- Input: {"registration_number": "MH12AB1234"}
- Output: {"vehicle_id": "V789012", "model": "Swift", ...}

POST /api/dealers/nearby
- Input: {"city": "Mumbai", "area": "Bandra"}
- Output: [{"dealer_id": "D345", "name": "Mumbai Maruti", ...}, ...]

# etc. - 11 total endpoints
```

### Mode 2: MSIL APIM (Production)

```
API_GATEWAY_MODE: str = "msil_apim"
MSIL_APIM_BASE_URL: str = "https://apim-dev.marutisuzuki.com"
MSIL_APIM_SUBSCRIPTION_KEY: Optional[str] = "subscription-key-xxxxxx"
MSIL_APIM_CLIENT_ID: Optional[str] = "client-id-xxxxxx"
MSIL_APIM_CLIENT_SECRET: Optional[str] = "client-secret-xxxxxx"
MSIL_APIM_TENANT_ID: Optional[str] = "tenant-id-xxxxxx"

# When tool is executed:
# Tool endpoint: /api/customer/resolve
# Actual URL: https://apim-dev.marutisuzuki.com/api/customer/resolve
# Headers include: Ocp-Apim-Subscription-Key

# Authentication:
# - OAuth2 client credentials flow
# - Token exchanged via Azure AD (OIDC)
```

## Request Flow by API Gateway Mode

Tool Executor (app/core/tools/executor.py)

```

_get_base_url():
  if API_GATEWAY_MODE == "mock":
    return "http://localhost:8080"
  else: # msil_apim
    return "https://apim-dev.marutisuzuki.com"

_get_headers(tool_auth_type):
  headers = {
    "Content-Type": "application/json",
    "Accept": "application/json"
  }

  if API_GATEWAY_MODE == "msil_apim":
    if MSIL_APIM_SUBSCRIPTION_KEY:
      headers["Ocp-Apim-Subscription-Key"] = MSIL_APIM_SUBSCRIPTION_KEY
      # Add OAuth token if available
    else:
      headers["X-API-Key"] = API_KEY

  return headers

execute(tool_name, arguments):
  base_url = _get_base_url()
  headers = _get_headers(tool.auth_type)

  url = f"{base_url}{tool.api_endpoint}"

  response = await http_client.post(
    url,
    headers=headers,
    json=arguments,
    timeout=30.0
  )

  return response.json()

```

## Security Policies & Rules

### RBAC Policies

**File:** mcp-server/app/core/policy/engine.py

```
# Simple RBAC rules (fallback if OPA not available):
```

```
SIMPLE_RULES = {
    "admin": [
        "*" # Wildcard: admin can do everything
    ],
    "developer": [
        "tools:read",
        "tools:write",
        "tools:create",
        "tools:delete",
        "policies:read",
        "specs:read"
    ],
    "operator": [
        "tools:read",
        "tools:invoke",
        "metrics:read",
        "audit:read_own"
    ],
    "user": [
        "tools:invoke",
        "tools:read_assigned"
    ]
}
```

## Risk-Based Policies

File: [mcp-server/app/core/policy/risk\\_policy.py](#)

```

RISK_POLICIES = {
    "read": RiskPolicy(
        risk_level="read",
        min_role="operator",           # operator, developer, admin can
access
        requires_elevation=False,
        requires_approval=False,
        rate_limit_tier="permissive",
        max_concurrency=50,
        pii_policy="redact"
    ),
    "write": RiskPolicy(
        risk_level="write",
        min_role="developer",         # developer, admin can access
        requires_elevation=False,
        requires_approval=False,
        rate_limit_tier="standard",
        max_concurrency=20,
        pii_policy="mask"
    ),
    "privileged": RiskPolicy(
        risk_level="privileged",
        min_role="admin",             # admin only, with elevation &
approval
        requires_elevation=True,       # Requires PIM approval
        requires_approval=True,        # Requires manager approval
        rate_limit_tier="strict",
        max_concurrency=5,
        pii_policy="redact_sensitive"
    )
}

```

## Rate Limit Policies

File: mcp-server/app/core/cache/rate\_limiter.py

```

RATE_LIMIT_TIERS = {
    "permissive": {
        "requests_per_minute": 100,
        "burst_size": 20,
        "window": 60,
        "retry_after": 60
    },
    "standard": {
        "requests_per_minute": 50,
        "burst_size": 10,
        "window": 60,
        "retry_after": 60
    },
    "strict": {
        "requests_per_minute": 10,
        "burst_size": 2,
        "window": 60,
        "retry_after": 60
    }
}

```

## RFP Requirements Mapping

### Complete Traceability Matrix

RFP Requirement	Implementation	File	Lines	Status
<b>Authentication: OAuth2/OIDC</b>	JWT token validation	<a href="#">oauth2_provider.py</a>	95-120	<input checked="" type="checkbox"/>
<b>Authorization: RBAC</b>	Role-based access control	<a href="#">policy/engine.py</a>	1-50	<input checked="" type="checkbox"/>
<b>Authorization: Tool-Level</b>	Per-tool access rules	<a href="#">policy/risk_policy.py</a>	39-90	<input checked="" type="checkbox"/>
<b>Policy Engine: OPA Integration</b>	OpenPolicy Agent ready	<a href="#">policy/engine.py</a>	1-30	<input checked="" type="checkbox"/>
<b>Rate Limiting: Per-User</b>	Token bucket by user_id	<a href="#">rate_limiter.py</a>	85-95	<input checked="" type="checkbox"/>
<b>Rate Limiting: Per-Tool</b>	Token bucket by tool_name	<a href="#">rate_limiter.py</a>	100-110	<input checked="" type="checkbox"/>
<b>Input Validation</b>	JSON schema validation	<a href="#">mcp.py</a>	200-250	<input checked="" type="checkbox"/>
<b>Audit Logging</b>	All actions logged	<a href="#">audit/service.py</a>	1-100	<input checked="" type="checkbox"/>

RFP Requirement	Implementation	File	Lines	Status
<b>PII Protection</b>	Masking + redaction	<code>audit/pii_masker.py</code>	1-50	✓
<b>Tool Discovery</b>	MCP tools/list method	<code>mcp.py</code>	100-150	✓
<b>Tool Execution</b>	MCP tools/call method	<code>mcp.py</code>	150-220	✓
<b>Idempotency</b>	For write operations	<code>idempotency/store.py</code>	1-80	✓
<b>Admin Tool Mgmt</b>	Zero-code OpenAPI import	<code>openapi_import.py</code>	1-200	✓
<b>Metrics</b>	Real-time analytics	<code>metrics/collector.py</code>	1-100	✓
<b>Multi-Gateway</b>	Mock API + MSIL APIM	<code>executor.py</code>	30-50	✓
<b>Error Handling</b>	JSON-RPC error codes	<code>mcp.py</code>	250-300	✓
<b>Correlation IDs</b>	Request tracing	<code>mcp.py</code>	60-70	✓
<b>Secrets Management</b>	Environment variables	<code>config.py</code>	1-100	✓

## Quick Reference: Key File Locations

### Authentication & Security

- JWT Handler: `mcp-server/app/core/auth/jwt_handler.py`
- OAuth2 Provider: `mcp-server/app/core/auth/oauth2_provider.py`
- Auth Models: `mcp-server/app/core/auth/models.py`

### Authorization & Policy

- Policy Engine: `mcp-server/app/core/policy/engine.py`
- Risk Policy: `mcp-server/app/core/policy/risk_policy.py`
- Policy Models: `mcp-server/app/core/policy/models.py`

### Rate Limiting & Caching

- Rate Limiter: `mcp-server/app/core/cache/rate_limiter.py`
- Cache Service: `mcp-server/app/core/cache/service.py`
- Idempotency: `mcp-server/app/core/idempotency/store.py`

### Tool Management

- Tool Registry: [mcp-server/app/core/tools/registry.py](#)
- Tool Executor: [mcp-server/app/core/tools/executor.py](#)
- OpenAPI Parser: [mcp-server/app/core/openapi/parser.py](#)

## Monitoring & Logging

- Metrics Collector: [mcp-server/app/core/metrics/collector.py](#)
- Audit Service: [mcp-server/app/core/audit/service.py](#)
- PII Masker: [mcp-server/app/core/audit/pii\\_masker.py](#)

## API Endpoints

- MCP Protocol: [mcp-server/app/api/mcp.py](#)
- Admin API: [mcp-server/app/api/admin.py](#)
- OpenAPI Import: [mcp-server/app/api/openapi\\_import.py](#)
- Analytics: [mcp-server/app/api/analytics.py](#)

## Configuration & Database

- Configuration: [mcp-server/app/config.py](#)
- Database: [mcp-server/app/db/database.py](#)
- Repositories: [mcp-server/app/db/repositories.py](#)

## UI Components

- Admin Portal: [admin-ui/src/pages/Import.tsx](#), [Dashboard.tsx](#), [Tools.tsx](#)
- Chat Interface: [chat-ui/src/App.tsx](#)

---

## Conclusion

This comprehensive documentation provides:

1. **Complete flow** from admin adding APIs to MCP client execution
2. **Multi-layered security** with authentication → authorization → rate limiting
3. **All data storage locations** with database schemas and file paths
4. **Role-based access control** with hierarchy and permissions
5. **Configuration management** across environments
6. **RFP requirements traceability** with file locations

All sensitive operations are logged, audited, and protected by role-based policies and rate limiting.

**Last Updated:** January 31, 2026